

## ▼ Install necessary libraries using pip

```
1 pip install numpy
```

```
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.23.5)
```

```
1 pip install pandas
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (1.5.3)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.3.post1)
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/dist-packages (from pandas) (1.23.5)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas) (1.16.0)
```

```
1 pip install scikit-learn
```

```
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.2.2)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.23.5)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.3)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.2.0)
```

```
1 pip install tensorflow
```

```
Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages (2.14.0)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=23.5.26 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (23.5.26)
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.5.4)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.9.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (16.0.6)
Requirement already satisfied: ml-dtypes==0.2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: numpy>=1.23.5 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.23.5)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.3.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow) (23.2)
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (67.7.2)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.3.0)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.5.0)
Requirement already satisfied: wrapt<1.15,>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.14.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.34.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.59.2)
Requirement already satisfied: tensorboard<2.15,>=2.14 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.14.1)
Requirement already satisfied: tensorflow-estimator<2.15,>=2.14.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.14.0)
Requirement already satisfied: keras<2.15,>=2.14.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.14.0)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.41.0)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.14.0)
Requirement already satisfied: google-auth-oauthlib<1.1,>=0.5 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.5.1)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.3)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.28.1)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.17.0)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.0.1)
Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (5.3.0)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.3.0)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.9)
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.3.1)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.2.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.0.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2023.7.22)
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.1.2)
Requirement already satisfied: pyasn1<0.6.0,>=0.4.6 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.5.0)
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.2.0)
```

## ▼ Import required libraries

```
1 import pandas as pd
2 import numpy as np
3 import tensorflow as tf
4 from tensorflow import keras
```

```

5 from keras.models import Sequential
6 from keras.layers import Dense
7 from sklearn.ensemble import RandomForestRegressor
8 from sklearn.metrics import r2_score
9 from sklearn.preprocessing import StandardScaler, MinMaxScaler
10 from sklearn.ensemble import RandomForestRegressor

```

## ▼ Read the dataset from a CSV file and display the first few rows

```
1 df = pd.read_csv('car_purchasing.csv',encoding='ISO-8859-1')
```

```
1 df.head()
```

	customer name	customer e-mail	country	gender	age
0	Martina Avila	cubilia.Curae.Phasellus@quisaccumsanconvallis.edu	Bulgaria	0	41.851720
1	Harlan Barnes	eu.dolor@diam.co.uk	Belize	0	40.870623
2	Naomi Rodriguez	vulputate.mauris.sagittis@ametconsectetueradip...	Algeria	1	43.152897
3	Jade Cunningham	malesuada@dignissim.com	Cook Islands	1	58.271369
4	Cedric Leach	felis.ullamcorper.viverra@egetmollislectus.net	Brazil	1	57.313749

## ▼ Define a list of columns to drop from the DataFrame

```

1 columns_to_drop = ['customer name', 'customer e-mail', 'country']
2 new_df = df.drop(columns=columns_to_drop)

```

ChatGPT

Show ChatGPT

```
1 new_df.head()
```

	gender	age	annual Salary	credit card debt	net worth	car purchase amount
0	0	41.851720	62812.09301	11609.380910	238961.2505	35321.45877
1	0	40.870623	66646.89292	9572.957136	530973.9078	45115.52566
2	1	43.152897	53798.55112	11160.355060	638467.1773	42925.70921
3	1	58.271369	79370.03798	14426.164850	548599.0524	67422.36313
4	1	57.313749	59729.15130	5358.712177	560304.0671	55915.46248

## ▼ Apply Min-Max scaling to the 'age' column in the DataFrame

```

1 #standard_scaler = StandardScaler()
2 #new_df['age_standard_scaled'] = standard_scaler.fit_transform(new_df[['age']])
3 min_max_scaler = MinMaxScaler()
4 new_df['age'] = min_max_scaler.fit_transform(new_df[['age']])

```

Show ChatGPT

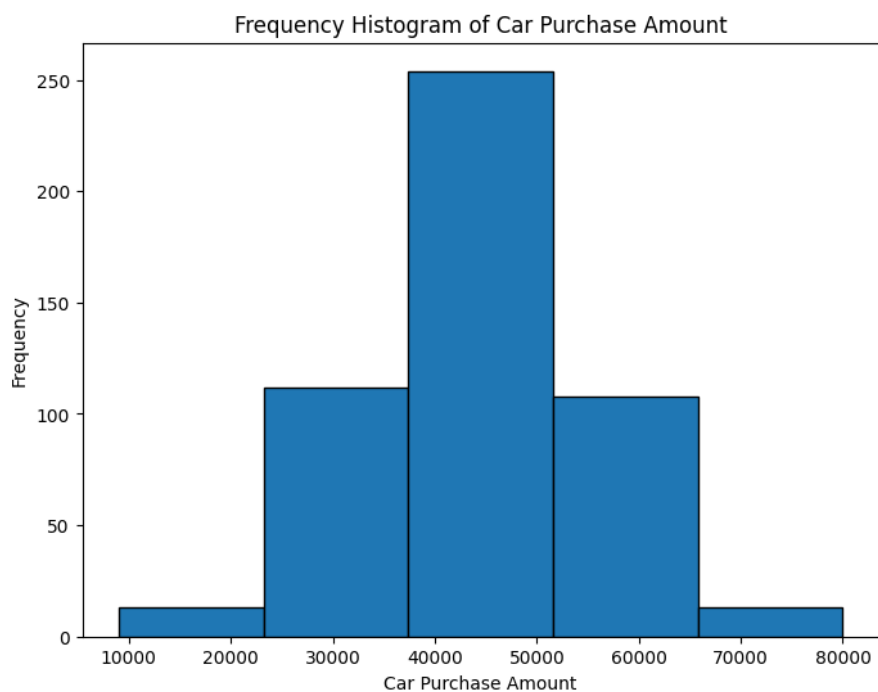
```
1 new_df.head()
```

	gender	age	annual Salary	credit card debt	net worth	car purchase amount	age_standard_s
0	0	41.851720	62812.09301	11609.380910	238961.2505	35321.45877	-0.5
1	0	40.870622	66646.80202	9572.057126	530072.0078	45115.52566	0.6

## ▼ Import Matplotlib for data visualization

```
1 import matplotlib.pyplot as plt
2 import pandas as pd
3
4 # Plot the frequency histogram
5 plt.figure(figsize=(8, 6))
6 plt.hist(df['car purchase amount'], bins=5, edgecolor='k')
7 plt.xlabel('Car Purchase Amount')
8 plt.ylabel('Frequency')
9 plt.title('Frequency Histogram of Car Purchase Amount')
10 plt.show()
```

Show ChatGPT



Create a frequency histogram for the 'car purchase amount' column

## ▼ Apply Standard Scaling to specific columns in the DataFrame

```
1 salary_scaler = StandardScaler()
2 debt_scaler = StandardScaler()
3 net_worth_scaler = StandardScaler()
4 new_df['annual Salary'] = salary_scaler.fit_transform(new_df[['annual Salary']])
5 new_df['credit card debt'] = debt_scaler.fit_transform(new_df[['credit card debt']])
6 new_df['net worth'] = net_worth_scaler.fit_transform(new_df[['net worth']])
7
8
```

Show ChatGPT

```
1 new_df.head()
```

	gender	age	annual Salary	credit card debt	net worth	car purchase amount	age_standard_scaled	age_min_max_scaled
0	0	0.437034	0.058576	0.574271	-1.110469	35321.45877	-0.550749	0.437034
1	0	0.417412	0.386570	-0.009951	0.573929	45115.52566	-0.673834	0.417412
2	1	0.463058	-0.712361	0.445452	1.193976	42925.70921	-0.387508	0.463058
3	1	0.765427	1.474794	1.382369	0.675595	67422.36313	1.509206	0.765427
4	1	0.746275	-0.205111	-1.218962	0.743113	55915.46248	1.389066	0.746275

+ Code + Text

```
1 new_df = new_df.iloc[:, :-2]
2
```

Show ChatGPT

```
1 new_df.head()
```

	gender	age	annual Salary	credit card debt	net worth	car purchase amount
0	0	0.437034	0.058576	0.574271	-1.110469	35321.45877
1	0	0.417412	0.386570	-0.009951	0.573929	45115.52566
2	1	0.463058	-0.712361	0.445452	1.193976	42925.70921
3	1	0.765427	1.474794	1.382369	0.675595	67422.36313
4	1	0.746275	-0.205111	-1.218962	0.743113	55915.46248

## ▼ Create a Random Forest Regressor model and fit it to the data

Calculate and print feature importances in descending order

```
1 X = new_df.drop('car purchase amount', axis=1)
2 y = new_df['car purchase amount']
3 model = RandomForestRegressor()
4 model.fit(X, y)
5 feature_importances = model.feature_importances_
6 sorted_indices = feature_importances.argsort()[::-1]
7 top_features = X.columns[sorted_indices[:3]]
8 print(top_features)
9
```

Index(['age', 'annual Salary', 'net worth'], dtype='object')

## ▼ Split the data into training and testing sets

```
1 from sklearn.model_selection import train_test_split
2 X = new_df[['age', 'annual Salary', 'net worth']]
3 y = new_df['car purchase amount']
4 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
5
```

Show ChatGPT

## ▼ Create a Sequential model for neural network

```
1 model = Sequential()
2 model.add(Dense(2, input_dim=3, activation='relu'))
3 model.add(Dense(1, activation='linear'))
```

## ▼ Compile the neural network model

```

1 model.compile(
2     loss = 'binary_crossentropy',
3     optimizer = 'adam',
4     metrics = ['accuracy']
5 )

```

## ▼ Fit the model using the original data

```

1 original_X = df[['age', 'annual Salary', 'net worth']]
2 original_y = df['car purchase amount']
3 model.fit(
4     original_X , original_y ,
5     epochs = 50 ,
6     batch_size = 3 ,
7     verbose = 1
8 )
9

```

167/167 [=====] - 0s 2ms/step - loss: -674150.4375 - accuracy: 0.0000e+00  
 Epoch 23/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.5625 - accuracy: 0.0000e+00  
 Epoch 24/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.5625 - accuracy: 0.0000e+00  
 Epoch 25/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.6250 - accuracy: 0.0000e+00  
 Epoch 26/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.2500 - accuracy: 0.0000e+00  
 Epoch 27/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.5625 - accuracy: 0.0000e+00  
 Epoch 28/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.3125 - accuracy: 0.0000e+00  
 Epoch 29/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.3750 - accuracy: 0.0000e+00  
 Epoch 30/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.3750 - accuracy: 0.0000e+00  
 Epoch 31/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.4375 - accuracy: 0.0000e+00  
 Epoch 32/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.5000 - accuracy: 0.0000e+00  
 Epoch 33/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.5000 - accuracy: 0.0000e+00  
 Epoch 34/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.6250 - accuracy: 0.0000e+00  
 Epoch 35/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.3125 - accuracy: 0.0000e+00  
 Epoch 36/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.6250 - accuracy: 0.0000e+00  
 Epoch 37/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.3750 - accuracy: 0.0000e+00  
 Epoch 38/50  
 167/167 [=====] - 1s 3ms/step - loss: -674150.3750 - accuracy: 0.0000e+00  
 Epoch 39/50  
 167/167 [=====] - 1s 3ms/step - loss: -674150.8125 - accuracy: 0.0000e+00  
 Epoch 40/50  
 167/167 [=====] - 1s 3ms/step - loss: -674150.5625 - accuracy: 0.0000e+00  
 Epoch 41/50  
 167/167 [=====] - 1s 3ms/step - loss: -674150.7500 - accuracy: 0.0000e+00  
 Epoch 42/50  
 167/167 [=====] - 1s 3ms/step - loss: -674150.5000 - accuracy: 0.0000e+00  
 Epoch 43/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.5625 - accuracy: 0.0000e+00  
 Epoch 44/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.3750 - accuracy: 0.0000e+00  
 Epoch 45/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.8750 - accuracy: 0.0000e+00  
 Epoch 46/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.8125 - accuracy: 0.0000e+00  
 Epoch 47/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.5000 - accuracy: 0.0000e+00  
 Epoch 48/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.5625 - accuracy: 0.0000e+00  
 Epoch 49/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.3125 - accuracy: 0.0000e+00  
 Epoch 50/50  
 167/167 [=====] - 0s 2ms/step - loss: -674150.6250 - accuracy: 0.0000e+00  
 <keras.src.callbacks.History at 0x7e0bdb354370>

## ▼ Fit the model using the training data

```

1 model.fit(
2     X_train , y_train ,
3     epochs = 50 ,
4     batch_size = 3 ,
5     verbose = 1
6 )

```

Epoch 1/50  
134/134 [=====] - 2s 7ms/step - loss: -670128.9375 - accuracy: 0.0000e+00  
Epoch 2/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8125 - accuracy: 0.0000e+00  
Epoch 3/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8750 - accuracy: 0.0000e+00  
Epoch 4/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.6875 - accuracy: 0.0000e+00  
Epoch 5/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.7500 - accuracy: 0.0000e+00  
Epoch 6/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.6875 - accuracy: 0.0000e+00  
Epoch 7/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8750 - accuracy: 0.0000e+00  
Epoch 8/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.7500 - accuracy: 0.0000e+00  
Epoch 9/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8750 - accuracy: 0.0000e+00  
Epoch 10/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.7500 - accuracy: 0.0000e+00  
Epoch 11/50  
134/134 [=====] - 0s 2ms/step - loss: -670129.0000 - accuracy: 0.0000e+00  
Epoch 12/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.9375 - accuracy: 0.0000e+00  
Epoch 13/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8750 - accuracy: 0.0000e+00  
Epoch 14/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.9375 - accuracy: 0.0000e+00  
Epoch 15/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8750 - accuracy: 0.0000e+00  
Epoch 16/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.9375 - accuracy: 0.0000e+00  
Epoch 17/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.7500 - accuracy: 0.0000e+00  
Epoch 18/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.7500 - accuracy: 0.0000e+00  
Epoch 19/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.9375 - accuracy: 0.0000e+00  
Epoch 20/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8125 - accuracy: 0.0000e+00  
Epoch 21/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8750 - accuracy: 0.0000e+00  
Epoch 22/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.7500 - accuracy: 0.0000e+00  
Epoch 23/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8125 - accuracy: 0.0000e+00  
Epoch 24/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8750 - accuracy: 0.0000e+00  
Epoch 25/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.7500 - accuracy: 0.0000e+00  
Epoch 26/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.7500 - accuracy: 0.0000e+00  
Epoch 27/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.8750 - accuracy: 0.0000e+00  
Epoch 28/50  
134/134 [=====] - 0s 2ms/step - loss: -670129.0000 - accuracy: 0.0000e+00  
Epoch 29/50  
134/134 [=====] - 0s 2ms/step - loss: -670128.9375 - accuracy: 0.0000e+00

## ▼ Evaluate the model on the testing data

```

1 model.evaluate(X_test,y_test,verbose=1)

```

4/4 [=====] - 0s 4ms/step - loss: -690237.1250 - accuracy: 0.0000e+00  
[-690237.125, 0.0]

## ▼ Create a Random Forest Regressor model with 100 estimators and fit it to the training data

```
1 random_forest = RandomForestRegressor(n_estimators=100)
2 random_forest.fit(X_train,y_train)
3
4
```

0.96004468314316

## ▼ Make predictions using the random forest model and calculate the R-squared score

```
1 random_f_pred = random_forest.predict(X_test)
2 r2_score(y_test,random_f_pred
```